

HedonicModel:SND

LuisM

3/22/2021

VARIABLES SELECTION

Loading the data

```
library(readxl)
df <- read_excel("/Users/Unimooc/Dropbox/2021/Directorio R/SpaceNextDoor/NextDoor/Optimal pricing/NextD
                sheet = "Data1")

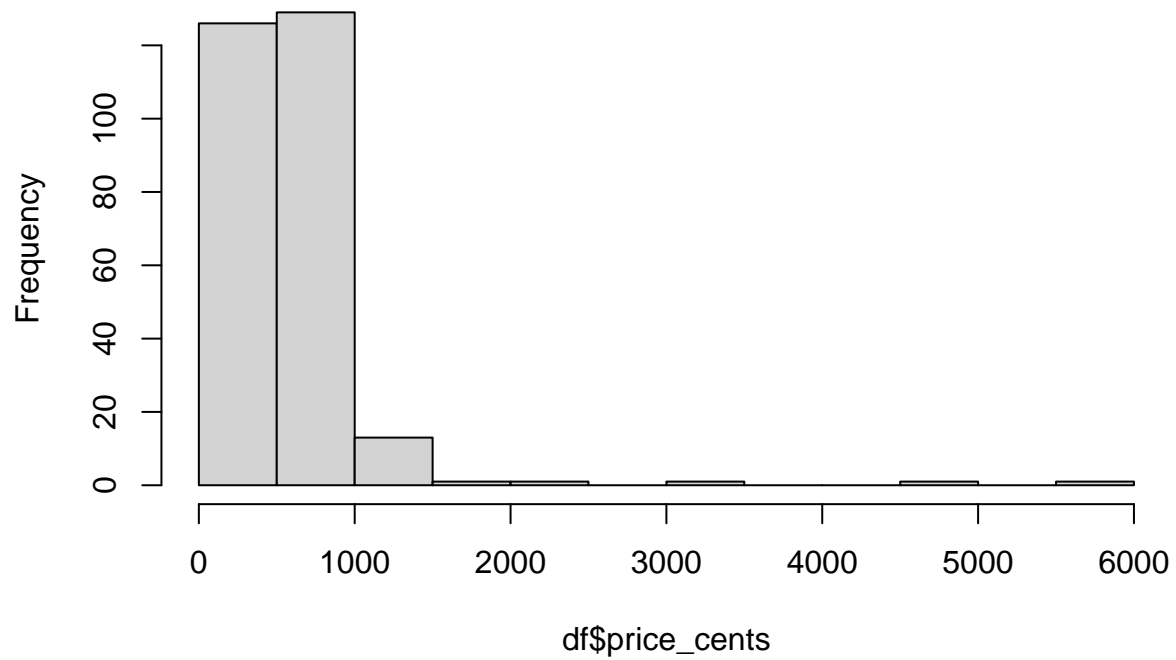
head(df)
```

Visualization

```
library(ggplot2)
library(corrplot)
library(tidyverse)
library(MASS)
```

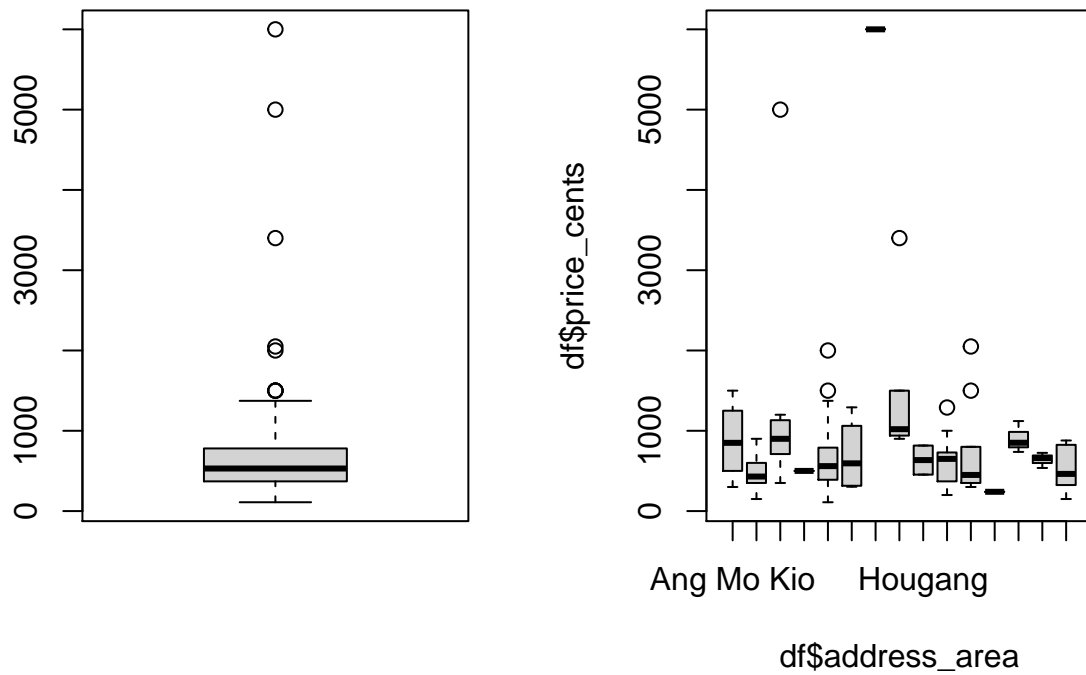
```
hist(df$price_cents)
```

Histogram of df\$price_cents



#The histogram shows a positive skweness (positioned to the left). There are many outliers that should


```
par(mfrow=c(1,2))  
  
boxplot(df$price_cents)  
boxplot(df$price_cents ~ df$address_area)
```



#The Boxplot confirms the presence of outliers caused by some isolated data in a few districts. As we g

```
df.cor <- cor(df[,c(21:42)], method = "kendall")
round(df.cor, digits = 1)
```

```
corrplot(df.cor, method = "shade",
          tl.col = "black",
          order = "AOE", type = "lower", diag = F)
```



#The correlation matrix indicates the presence of strong autocorrelation between some variables. We sho

Training and test sample division

```
library(dplyr)
library(caret)
```

```
df.sel <- df[,-c(1,3,4,6:20)]

set.seed(2021)

dfPartition <- createDataPartition(y = df.sel$price_cents,
                                     p = 0.7, list = F)

Training <- df.sel[dfPartition,]
```

```
## Warning: The 'i' argument of '['()' can't be a matrix as of tibble 3.0.0.
## Convert to a vector.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_warnings()' to see where this warning was generated.
```

```
Test <- df.sel[-dfPartition,]
```

Variables selection

AIC forward selection

```
Modelzero <- lm(price_cents~1,data=Training)
summary(Modelzero)

FitAll = lm(price_cents ~ ., data=Training)
formula(FitAll)

model.forward <- step(Modelzero,direction="forward",scope=formula(FitAll))
```

```
summary(model.forward)
```

```
##
## Call:
## lm(formula = price_cents ~ 'area squared foot' + 'Lift access' +
##     'Step ladder' + Locker + CCTV + '24 hours' + 'Loading dock' +
##     'Free Parking' + WIFI + 'On-Site staff' + 'Entire room' +
##     'Lockable doors' + Unit + Sprinkler + Security + 'Wine cellar' +
##     'Part room', data = Training)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -473.97  -72.02   -7.73   70.93  515.74
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      985.189    146.850   6.709 2.62e-10 ***
## 'area squared foot'    55.976     2.058  27.201 < 2e-16 ***
## 'Lift access'      -799.348    100.674  -7.940 2.35e-13 ***
## 'Step ladder'       140.546     96.978   1.449 0.149055
## Locker             -59.443     82.170  -0.723 0.470388
## CCTV             -325.207    180.135  -1.805 0.072739 .
## '24 hours'         493.892    163.015   3.030 0.002819 **
## 'Loading dock'      131.575     25.772   5.105 8.56e-07 ***
## 'Free Parking'     -91.800     25.972  -3.535 0.000523 ***
## WIFI              -17.075     83.769  -0.204 0.838718
## 'On-Site staff'     152.321     99.474   1.531 0.127510
## 'Entire room'     -236.118     80.769  -2.923 0.003920 **
## 'Lockable doors'    152.104     66.511   2.287 0.023400 *
## Unit              274.955     86.167   3.191 0.001681 **
## Sprinkler          130.131     49.903   2.608 0.009903 **
## Security          -595.850    167.229  -3.563 0.000472 ***
## 'Wine cellar'       219.157     91.442   2.397 0.017598 *
## 'Part room'       -219.952     96.771  -2.273 0.024248 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 137.4 on 175 degrees of freedom
## Multiple R-squared:  0.9214, Adjusted R-squared:  0.9138
## F-statistic: 120.7 on 17 and 175 DF,  p-value: < 2.2e-16
```

Forward model:

```
price_cents ~ 'area squared foot' + 'Lift access' + 'Step ladder' + Locker + CCTV + '24 hours' + 'Loading dock'
```

```
predict.for.tr <- predict(model.forward, newdata = Training)

training.for.mse <- mean((predict.for.tr - Training$price_cents)^2)
paste("Training MSE error:", training.for.mse)
```

```
## [1] "Training MSE error: 17119.3585103714"
```

```
predict.for.tst <- predict(model.forward, newdata = Test)

test.for.mse <- mean((predict.for.tst - Test$price_cents)^2)
paste("Test MSE error:", test.for.mse)
```

```
## [1] "Test MSE error: 415545.240857092"
```

Summary

```
price_cents ~ 'area squared foot' + 'Lift access' + 'Step ladder' + Locker + CCTV + '24 hours' + 'Loading dock'
```

AIC backward selection

```
model.backward <- stepAIC(FitAll, trace=TRUE, direction="backward")
```

```
summary(model.backward)
```

```
##
## Call:
## lm(formula = price_cents ~ 'area squared foot' + 'Part room' +
##     'Entire room' + Unit + 'Wine cellar' + '24 hours' + 'Lift access' +
##     Security + 'Loading dock' + CCTV + 'Lockable doors' + 'Free Parking' +
##     Trolleys + 'Step ladder' + Sprinkler, data = Training)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -457.06  -71.35   -7.86    69.29   527.52
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      919.675    133.102   6.910 8.45e-11 ***
## 'area squared foot'    55.888     2.014  27.751 < 2e-16 ***
## 'Part room'      -231.221    86.147  -2.684 0.007964 **
## 'Entire room'     -235.911    78.712  -2.997 0.003117 **
## Unit              285.802    79.996   3.573 0.000455 ***
## 'Wine cellar'      292.047    53.050   5.505 1.28e-07 ***
## '24 hours'        491.699   153.464   3.204 0.001608 **
## 'Lift access'     -735.360    93.231  -7.888 3.06e-13 ***
```

```
## Security          -657.253    136.807  -4.804 3.30e-06 ***
## 'Loading dock'    131.412     25.154   5.224 4.88e-07 ***
## CCTV            -322.249    165.985  -1.941 0.053793 .
## 'Lockable doors'  140.312     64.461   2.177 0.030827 *
## 'Free Parking'    -90.558     25.458  -3.557 0.000481 ***
## Trolleys          136.907     88.730   1.543 0.124626
## 'Step ladder'     128.258     83.634   1.534 0.126923
## Sprinkler         130.705     48.089   2.718 0.007222 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 136.8 on 177 degrees of freedom
## Multiple R-squared:  0.9212, Adjusted R-squared:  0.9145
## F-statistic: 137.9 on 15 and 177 DF, p-value: < 2.2e-16
```

Summary

```
price_cents ~ 'area squared foot' + 'Part room' + 'Entire room' + Unit + 'Wine cellar' + '24 hours' + '1
```

```
predict.bck.tr <- predict(model.backward, newdata = Training)

training.bck.mse <- mean((predict.bck.tr - Training$price_cents)^2)
paste("Training MSE error:", training.bck.mse)
```

```
## [1] "Training MSE error: 17167.6882295733"
```

```
predict.bck.tst <- predict(model.backward, newdata = Test)

test.bck.mse <- mean((predict.bck.tst - Test$price_cents)^2)
paste("Test MSE error:", test.bck.mse)
```

```
## [1] "Test MSE error: 440903.002491439"
```

Ridge and Lasso regularizations

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'Matrix'
```

```
## The following objects are masked from 'package:tidyr':
```

```
##
```

```
## expand, pack, unpack
```


```
## Loaded glmnet 4.0-2
```

```
# Convert into a matrix train and test data
train.mat <- model.matrix(price_cents ~ ., data = Training)
test.mat <- model.matrix(price_cents ~ ., data = Test)
```

Ridge

```
# Cross validation to obtain the best value of lambda. Error evolution.
cv.ridge <- cv.glmnet(x = train.mat, y = Training$price_cents, alpha = 0,
                      lambda = NULL, type.measure="mse")

plot(cv.ridge)
```



```
paste("Best lambda:", cv.ridge$lambda.min)
```

```
## [1] "Best lambda: 91.8822861493194"
```

```
paste("Best lambda + y sd:", cv.ridge$lambda.1se)
```

```
## [1] "Best lambda + y sd: 370.930374462024"
```



```
# Training the model
mod.ridge.train <- glmnet(x = train.mat, y = Training$price_cents, alpha = 0,
                          lambda = cv.ridge$lambda.1se)
```

```
dim(coef(mod.ridge.train))
```

```
## [1] 25 1
```

```
coef(mod.ridge.train, s = "lambda.1se")
```

```
## 25 x 1 sparse Matrix of class "dgCMatrix"
##                               1
## (Intercept)                1050.104172
## (Intercept)                  .
## 'area squared foot'         29.275100
## Locker                      -103.832446
## Cupboard                     .
## 'Racking shelf'             -30.279797
## 'Part room'                 -84.405881
## 'Entire room'               73.527024
## Unit                       112.990315
## 'Wine cellar'              131.036923
## 'Air con'                   39.245186
## '24 hours'                 174.035224
## 'Lift access'              -352.066609
## Security                   -239.700388
## 'Loading dock'             70.563736
## CCTV                     -434.236379
## 'Lockable doors'          110.067463
## 'On-Site staff'            5.434037
## 'Free Parking'            -24.292767
## Trolleys                   10.110869
## 'Step ladder'              52.541545
## 'Racking/Shelving'        -30.444043
## Sprinkler                  -38.295701
## 'Power supply'            -18.012737
## WIFI                       14.963850
```

```
# Training predictions
pred.ridge <- predict(mod.ridge.train, newx = train.mat)
```

```
# Training error (MSE)
tr.ridge.mse <- mean((pred.ridge - Training$price_cents)^2)
paste("Training MSE error:", tr.ridge.mse)
```

```
## [1] "Training MSE error: 50145.6072531696"
```

```
#Test predictions: using training model
pred.test.ridge <- predict(mod.ridge.train, newx = test.mat)

test.ridge.mse <- mean((pred.test.ridge - Test$price_cents)^2)
paste("Test MSE error:", test.ridge.mse)
```


```
## [1] "Test MSE error: 413568.619418348"
```

Ridge model:

```
price_cents ~ 'area squared foot' + Locker + 'Racking shelf' + 'Part room'  
+ 'Entire room' + Unit + 'Wine cellar' + 'Air con' + '24 hours' + 'Lift access'  
+ Security + 'Loading dock' + CCTV + 'Lockable doors' + 'On-Site staff' + 'Free Parking'  
+ Trolleys + 'Step ladder' + 'Racking/Shelving' + Sprinkler + 'Power supply' + WIFI
```

Lasso

```
cv.lasso <- cv.glmnet(x = train.mat, y = Training$price_cents, alpha = 1,  
                      lambda = NULL, type.measure="mse")  
  
plot(cv.lasso)
```



```
paste("Best lambda:", cv.lasso$lambda.min)
```

```
## [1] "Best lambda: 26.7841738243326"
```

```
paste("Best lambda + y sd:", cv.lasso$lambda.1se)
```

```
## [1] "Best lambda + y sd: 98.5223630810571"
```

```
# Training the model
```

```
mod.lasso.train <- glmnet(x = train.mat, y = Training$price_cents, alpha = 1,  
                          lambda = cv.lasso$lambda.1se)
```

```
dim(coef(mod.lasso.train))
```

```
## [1] 25 1
```

```
coef(mod.lasso.train, s = "lambda.1se")
```

```
## 25 x 1 sparse Matrix of class "dgCMatrix"
```

```
##              1  
## (Intercept)    645.44705  
## (Intercept)      .  
## 'area squared foot' 43.68657  
## Locker            .  
## Cupboard          .  
## 'Racking shelf'    .  
## 'Part room'        .  
## 'Entire room'      .  
## Unit              45.59185  
## 'Wine cellar'      .  
## 'Air con'          .  
## '24 hours'         .  
## 'Lift access'      -194.63918  
## Security           .  
## 'Loading dock'     .  
## CCTV             -89.24854  
## 'Lockable doors'   .  
## 'On-Site staff'    .  
## 'Free Parking'     .  
## Trolleys           .  
## 'Step ladder'      .  
## 'Racking/Shelving' .  
## Sprinkler          .  
## 'Power supply'     .  
## WIFI              .
```

```
# Training predictions
```

```
pred.lasso <- predict(mod.lasso.train, newx = train.mat)
```

```
# Training error (MSE)
```

```
tr.lasso.mse <- mean((pred.lasso - Training$price_cents)^2)  
paste("Training MSE error:", tr.lasso.mse)
```

```
## [1] "Training MSE error: 57760.6240008345"
```

```
#Test predictions: using training model
pred.test.lasso <- predict(mod.lasso.train, newx = test.mat)

test.lasso.mse <- mean((pred.test.lasso - Test$price_cents)^2)
paste("Test MSE error:",test.lasso.mse)
```

```
## [1] "Test MSE error: 380197.505939534"
```

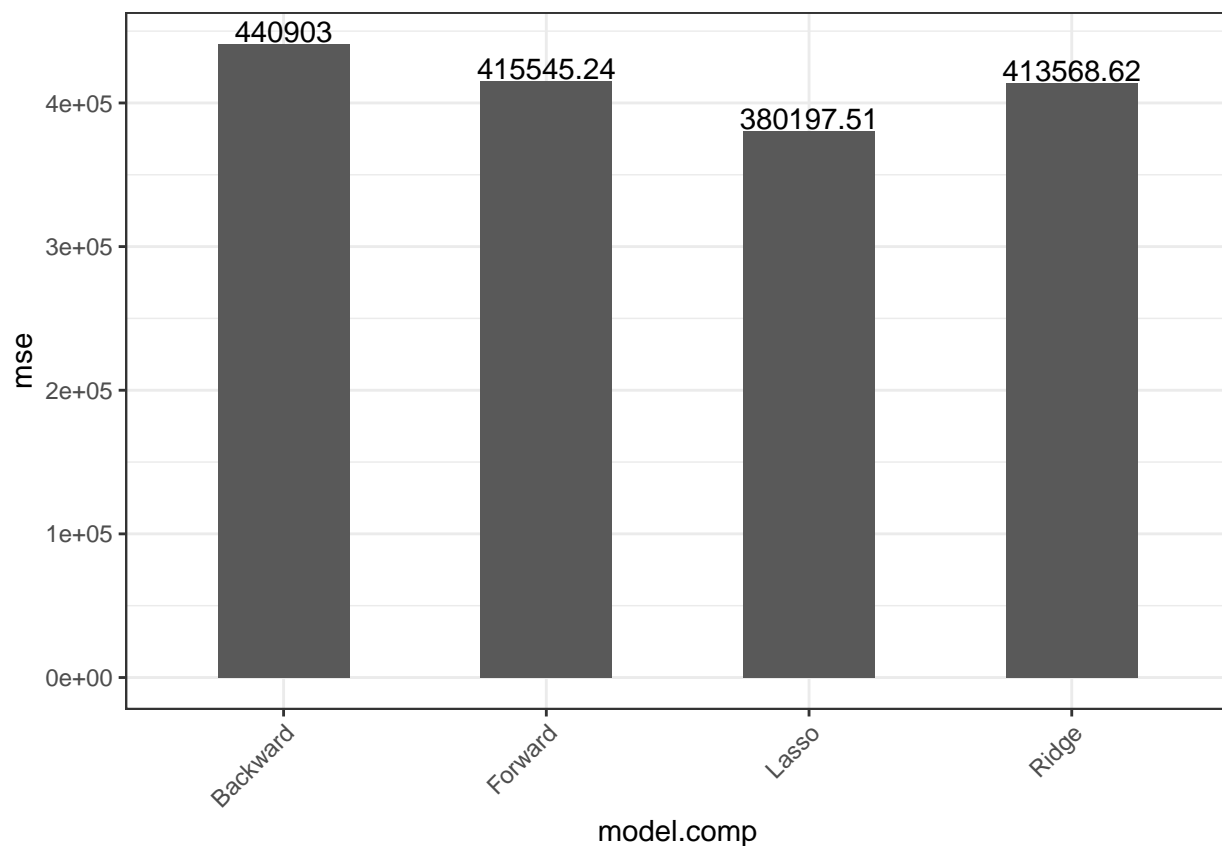
Lasso model:

```
price_cents ~ 'area squared foot' + Unit + 'Lift access' + CCTV
```

Comparing results

```
df_compar <- data.frame(model.comp = c("Forward", "Backward", "Ridge", "Lasso"),
mse = c(test.for.mse, test.bck.mse, test.ridge.mse, test.lasso.mse))

ggplot(data = df_compar, aes(x = model.comp, y = mse)) + geom_col(width = 0.5) +
  geom_text(aes(label = round(mse, 2)), vjust = -0.1) + theme_bw() + theme(axis.text.x = element_text(
hjust = 1))
```



Linear Model

```
lmodel.Train <- lm (price_cents ~ 'area squared foot' + Unit + 'Lift access' + CCTV, data = Training)

summary(lmodel.Train)
```

```
##
## Call:
## lm(formula = price_cents ~ 'area squared foot' + Unit + 'Lift access' +
##     CCTV, data = Training)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -581.08  -74.27  -20.80   58.26 1337.59
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1163.617    116.902   9.954 < 2e-16 ***
## 'area squared foot'    55.419     2.214  25.029 < 2e-16 ***
## Unit              233.773    44.983   5.197 5.25e-07 ***
## 'Lift access'      -670.286    106.657  -6.285 2.24e-09 ***
## CCTV             -235.437    93.576  -2.516  0.0127 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 189.3 on 188 degrees of freedom
## Multiple R-squared:  0.8398, Adjusted R-squared:  0.8364
## F-statistic: 246.3 on 4 and 188 DF,  p-value: < 2.2e-16
```

```
lmodel.Test <- lm (price_cents ~ 'area squared foot' + Unit + 'Lift access' + CCTV, data = Test)

summary(lmodel.Test)
```

```
##
## Call:
## lm(formula = price_cents ~ 'area squared foot' + Unit + 'Lift access' +
##     CCTV, data = Test)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -747.0 -149.4   30.5   69.4 4429.6
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -160.11    319.12  -0.502 0.617306
## 'area squared foot'    82.02    16.04   5.112 2.31e-06 ***
## Unit             643.57    159.82   4.027 0.000133 ***
## 'Lift access'         NA         NA      NA      NA
## CCTV           266.82    299.15   0.892 0.375248
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 554.9 on 76 degrees of freedom
## Multiple R-squared:  0.3849, Adjusted R-squared:  0.3606
## F-statistic: 15.85 on 3 and 76 DF,  p-value: 4.24e-08
```